


Machine Learning (ML):

- Subset of AI **(without explicit instructions)**
- Make predictions or decisions based on the patterns it's learned.

AWS AI/ML solutions

The AWS AI/ML stack is composed of the following three tiers of solutions:

- *AI services* - pre-built models that are already trained to perform specific functions
- *ML services* - a more customized approach with Amazon SageMaker AI where you build, train, and deploy your own ML models with fully managed infrastructure
- *ML frameworks and infrastructure* - a completely custom approach to building models using purpose-built chips that integrate with popular ML frameworks



Tier 1: Pre-built AWS AI services: **Managed******

- Pre-built models that are already trained to perform specific functions.

Language services:

Amazon Comprehend:

- Uses NLP to **extract key insights** from documents.
- Recognizing key phrases, language, sentiment.
- Use cases: Content classification, customer sentiment analysis, and compliance monitoring

Amazon Polly:

- Converts **text into speech**.
- Use cases: Virtual assistants, e-learning applications, and accessibility enhancements

Amazon Transcribe:

- Converts **speech into text**.
- Speaker identification, custom vocabulary, and real-time transcription.
- Use Cases: Automated subtitling, and metadata generation for media content.

Amazon Translate: - Text translation service.

Computer vision & search services:

Amazon Kendra:

- Uses **NLP to search for answers within large enterprise content.**
- Understands context of a query; can return more precise and relevant answers than just a list of documents with matching keywords.
- Use cases: Intelligent search, chatbots, and application search integration

Amazon Rekognition:

- **Video analysis service.**
- Identify objects, people, text, scenes, and activities within images and videos stored in S3
- Content moderation, identity verification, media analysis, and home automation experiences

Amazon Textract:

- Detects & extracts typed/handwritten text found in documents
- Use cases: Financial, healthcare, and government form text extraction for quick processing

Conversational AI personalization services:

Amazon Lex:

- Add **voice & text conversational interfaces to applications.**
- Uses both NLU and automatic speech recognition (ASR) to create lifelike conversations.
- Use cases: Virtual assistants, natural language search for FAQs, and automated application bots

Amazon Personalize

- Use historical data to build intelligent applications with **personalized recommendations** for customers.
- Use cases: Personalized streaming, product, and trending recommendations

#ML frameworks:

- Software library/tool that provides experienced ML practitioners with **pre-built**, optimized components for building machine learning models.
- PyTorch, Apache M-X Net, and TensorFlow.

Generative AI on AWS:

Amazon SageMaker JumpStart:

- An ML hub with FMs and pre-built ML solutions deployable with a few clicks
- Rapid ML model deployments, Custom fine-tuned solutions, ML experiments and prototypes

Amazon Bedrock: ***Managed***

- Adapt & deploy FMs from Amazon and leading AI companies (Claude)
- Enterprise-grade generative AI, **Multimodal content generation**, **Advanced conversational AI**

Amazon Q:

- **Interactive AI assistant** that can be integrated with a company's information repositories

Amazon Q Business:

- Answer pressing questions, help solve problems, and take actions using the data and expertise found in company's information repositories.

Amazon Q Developer:

- Code recommendations to accelerate development for coding languages including C#, Java, JavaScript, Python, and TypeScript applications.
- Integrates with multiple IDEs.

AI/ML & Generative AI Services (Exam Cues)

Service	Key Purpose
Amazon SageMaker JumpStart	Pre-built ML models & solutions, fine-tuning, rapid deployment
Amazon Bedrock	Fully managed, multi-model generative AI (text, images, etc.) via single API
Amazon Q Business	AI assistant for org data – answers, workflows, insights
Amazon Q Developer	AI code generator & reviewer inside IDEs

AWS Data Pipeline Services (I-S-C-P-A):

1. Data Ingestion

Service	Use
Amazon Kinesis Data Streams	Real-time ingestion, massive throughput
Amazon Data Firehose	Near real-time ingestion to S3, Redshift, etc.

2. Data Storage

Service	Use
Amazon S3	Data lake for raw, structured, unstructured
Amazon Redshift	Data warehouse for structured / semi structured

3. Data Cataloging

Service	Use
AWS Glue Data Catalog	Central metadata repository

4. Data Processing

Service	Use
AWS Glue	Fully managed ETL
Amazon EMR	Big data frameworks (Spark, Hadoop, Hive)

5. Data Analysis & Visualization

Service	Use
Amazon Athena	Serverless SQL on S3 and other sources
Amazon Redshift	Frequent high-performance queries
Amazon QuickSight	Interactive dashboards
Amazon OpenSearch Service	Search & real-time monitoring (Visualize data)

